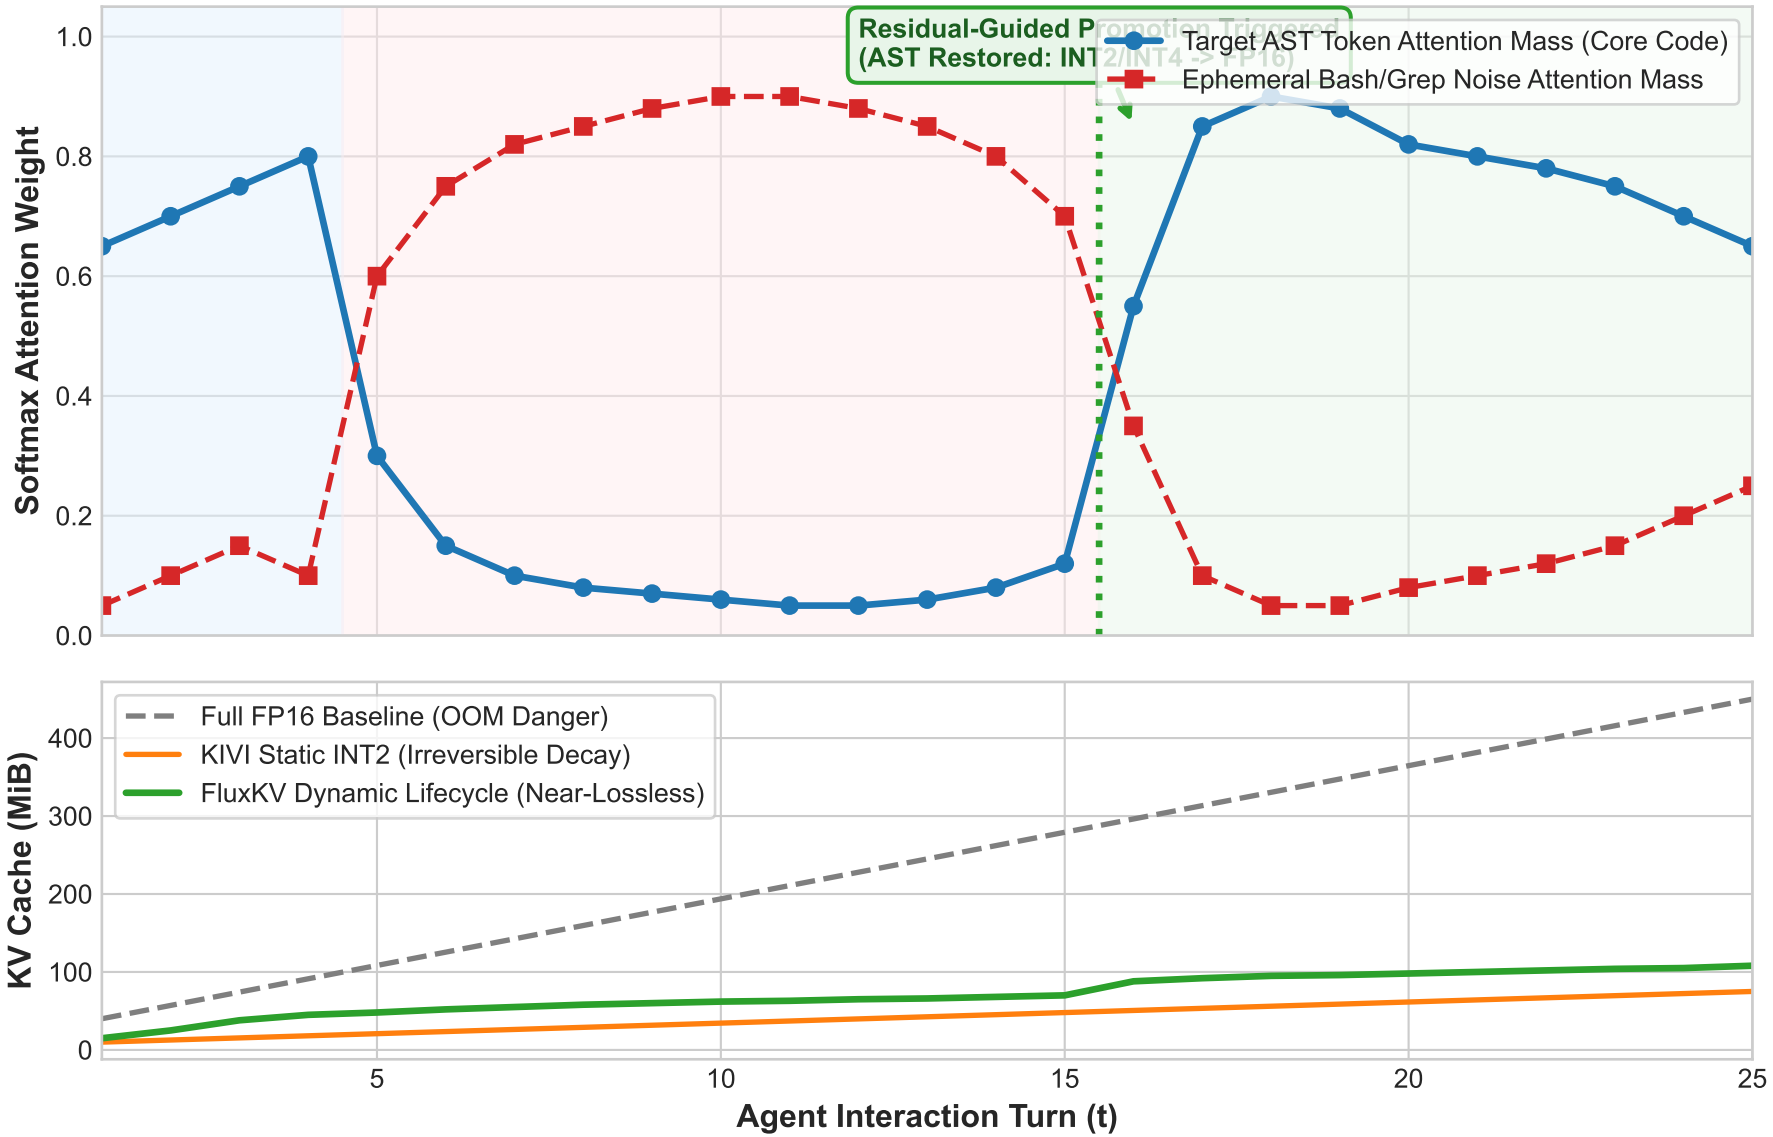
Multi-Turn Attention Dynamics and FluxKV Reversible Precision Promotion
Residual-Guided P
(AST Restored: INT
Target AST Token Attention Mass (Core Code)
Ephemeral Bash/Grep Noise Attention Mass
Softmax Attention Weight
1.0
0.8
0.6
0.4
0.2
0.0
Full FP16 Baseline (OOM Danger)
KIVI Static INT2 (Irreversible Decay)
FluxKV Dynamic Lifecycle (Near-Lossless)
KV Cache (MiB)
400
300
200
100
0
5
10
15
20
25
Agent Interaction Turn (t)